

As a TA for one of the large undergraduate classes at UC Berkeley (some approaching a thousand students), Rohan Lageweg found himself spending countless hours answering hundreds of emails from students. As the semester dragged on, he began to notice that many of the emails were similar in content, and many of his replies were almost exact replicas of each other. He began to wonder if there was a way to make this process more efficient.

To develop a solution, Rohan and five other UC Berkeley students used natural language processing and classification algorithms to come up with Flo, an email management system. This system takes the form of a Google Chrome extension, grouping gmail emails into useful categories and offering pre-written responses for each classification. The project involved downloading all course emails from that semester (with permission from course staff) to use as a dataset, anonymizing, cleaning and labeling the dataset, training various classification algorithms on the dataset and comparing their performances, and building a chrome extension and API to host the model.

